

# Conditioning and Directionality in Adversarial Transfer on CIFAR-10

Anonymous authors

Paper under double-blind review

## Abstract

Conventionally reported source attack success rate (ASR) and pairwise transfer rate (PTR) cannot reconstruct transfer conditional on source success (CTR): source ASR uses a broader denominator than the pair-dependent population required by CTR, and it omits target behavior when the source attack fails. We demonstrate this distinction on the full CIFAR-10 test set using three independently trained ResNet-18, VGG16, and MobileNetV2 checkpoints. The corrected evaluation excludes existing clean errors and reports exact denominators. At 2/255, conditioning is nearly mechanical in five of six directions: target error among eligible source failures is at most 0.68%, and mean CTR is within 0.27 points of the zero-residual value  $\text{PTR}/a_{st}$ . VGG16-to-MobileNetV2 is the exception, with a 2.83% failed-source residual and a 1.19-point difference. At 8/255, mean PGD PTR ranges from 18.09% for ResNet-to-VGG to 95.23% for VGG-to-MobileNet. A seed-0 control training all architectures at learning rate 0.01 preserves the VGG transfer asymmetry: its outbound PTR is 93.12–95.98%, versus 18.55–32.71% in reverse. The result is a seed-replicated, denominator-complete account of when source conditioning adds information and when it acts mainly as arithmetic rescaling.

## 1 Introduction

Adversarial examples expose a gap between ordinary predictive accuracy and behavior under deliberately chosen perturbations. FGSM provides a gradient-based construction (Goodfellow et al., 2015), while iterative optimization enables broader searches within a threat model (Madry et al., 2018; Carlini & Wagner, 2017). Transfer attacks use a source model to construct an input that is subsequently evaluated on another model. The target supplies neither gradients nor feedback during construction.

Transferability is an empirical property of a source, a target, an attack, and an evaluation population. It cannot be inferred from a high error rate alone. If the target already misclassifies most clean inputs, an unchanged input may appear to be a successful attack under an unconditional metric. A mismatch between label spaces makes this problem more severe.

The study asks whether conclusions about attack rankings and model-pair differences are stable under explicitly stated conditioning populations. Its contributions are:

1. an exact decomposition showing why conventional source ASR and PTR do not determine CTR;
2. a full-test, three-seed measurement with jointly clean populations, exact counts, and PGD PTR spanning 18.09–95.23% at 8/255; and
3. convergence and matched-learning-rate controls that preserve the main directional result while exposing its dependence on optimization.

The self-correction in Section 3 motivates the measurement question and is not evidence of a general failure rate across the literature.

## 2 Background and scope

Goodfellow et al. (2015) connect adversarial vulnerability to high-dimensional linear behavior and introduce FGSM. Madry et al. (2018) formulate robustness using optimization over a perturbation set. Carlini & Wagner (2017) develop optimization-based attacks and demonstrate the importance of careful evaluation. Papernot et al. (2017) demonstrate practical transfer-based black-box attacks using substitute models. Liu et al. (2017) study non-targeted and targeted transfer across large models, while Tramèr et al. (2017) characterize shared adversarial subspaces and limits of transfer.

Directionality and the tension between white-box optimization and transfer are established phenomena: Demontis et al. (2019) explain asymmetric transfer through target vulnerability and surrogate complexity, while Dong et al. (2018) use momentum to improve black-box transfer. We position the present study as a denominator-complete remeasurement of these phenomena rather than a new transfer attack.

Carlini et al. (2019) systematize the methodological foundations and common pitfalls of adversarial-robustness evaluation. Croce & Hein (2020) introduce AutoAttack, an ensemble intended to reduce dependence on attack tuning and improve evaluation reliability. Tramèr et al. (2020) emphasize defense-specific adaptive evaluation and illustrate how apparently adequate evaluations can miss weaknesses. RobustBench standardizes robustness benchmarking under specified tasks and threat models and uses AutoAttack as part of its evaluation framework (Croce et al., 2021).

Our study complements these evaluation efforts by examining source-to-target transfer and the populations on which transfer rates are conditioned. It does not introduce a stronger attack, certify robustness, or replace adaptive defense evaluation. The FGSM/PGD study does not constitute an AutoAttack evaluation; citing AutoAttack and RobustBench must not be interpreted as having run either benchmark. Attack-strength sensitivity and the source white-box controls are necessary for interpreting the transfer measurements.

The present protocol concerns untargeted digital attacks against deterministic image classifiers. It does not evaluate physical attacks, detection systems, adversarially trained defenses, or deployment security. CIFAR-10 results would not by themselves justify claims about autonomous vehicles or face recognition. The selected architectures are established CNN baselines, not a comprehensive representation of contemporary vision systems.

## 3 Motivating case study: a self-correction

An earlier implementation compared CIFAR-10 labels directly with the indices of 1,000-class ImageNet models, omitted required preprocessing, mixed perturbation coordinate systems, and counted existing clean errors as attack successes. It also lacked compatible checkpoints and per-example logs. Those rates and rankings are excluded from the revised empirical record. This self-correction motivates compatible labels, raw-pixel budgets, clean eligibility, and recorded source success; it is not the empirical contribution. The line-by-line audit is retained in the anonymized supplementary artifact.

## 4 Why error is not attack success

Let  $f_t$  be the target predictor,  $y$  the true label, and  $A_s(x)$  an attack constructed using source  $s$ . Define clean error  $e_t = \Pr[f_t(x) \neq y]$ , newly induced error  $q = \Pr[f_t(A_s(x)) \neq y \mid f_t(x) = y]$ , and persistence of existing error  $r = \Pr[f_t(A_s(x)) \neq y \mid f_t(x) \neq y]$ . By partitioning on clean correctness,

$$\Pr[f_t(A_s(x)) \neq y] = (1 - e_t)q + e_tr. \quad (1)$$

For the identity attack,  $q = 0$  and  $r = 1$ , so adversarial error equals clean error. An illustrative classifier with 99% clean error therefore has 99% adversarial error under an identity attack and zero newly induced errors. This is an algebraic example, not a measurement from CIFAR-10.

The unconditional attacked-error rate and CTR condition on different populations, so their difference is not an estimate of inflation caused by clean errors. Equation (1) attributes all-input attacked error to

newly induced and persistent clean errors; PTR and CTR then add source-correctness and source-success conditions, with empty conditioning populations remaining undefined.

#### 4.1 Recorded estimands

For a fixed evaluation set of  $N$  examples, let  $C_s(i)$  and  $C_t(i)$  indicate clean correctness. Let  $S_s(i)$  and  $S_t(i)$  indicate misclassification of the source-generated attacked input. Define the pairwise eligible set  $E_{st} = \{i : C_s(i) \wedge C_t(i)\}$  and successful-source subset  $F_{st} = \{i \in E_{st} : S_s(i)\}$ . We report

$$\text{ASR}_s = \frac{\sum_{i:C_s(i)} S_s(i)}{\sum_i C_s(i)}, \quad (2)$$

$$\text{PTR}_{st} = \frac{\sum_{i \in E_{st}} S_t(i)}{|E_{st}|}, \quad (3)$$

$$\text{CTR}_{st} = \frac{\sum_{i \in F_{st}} S_t(i)}{|F_{st}|}. \quad (4)$$

PTR measures target failure on jointly clean-correct inputs whether or not the source attack succeeded. CTR measures transfer conditional on source success. To expose the arithmetic effect of conditioning, define pairwise source success  $a_{st} = \Pr[S_s \mid E_{st}]$  and failed-source target error  $b_{st} = \Pr[S_t \mid E_{st}, \neg S_s]$ . Partitioning the eligible set gives the exact identity

$$\text{PTR}_{st} = a_{st} \text{CTR}_{st} + (1 - a_{st})b_{st}. \quad (5)$$

Consequently,  $\text{CTR}_{st} = \text{PTR}_{st} / a_{st}$  only when  $b_{st} = 0$ . The conventional  $\text{ASR}_s$  above cannot be substituted for  $a_{st}$  because its denominator includes source-correct inputs on which the target was already wrong. We report  $a_{st}$  and  $b_{st}$  for the conditioning analysis. The diagonal is a white-box control; its CTR is tautologically one whenever its denominator is positive.

The artifact tables carry numerator and denominator counts for every reported fraction. An empty denominator is undefined, not zero. Clean accuracy and attacked accuracy on the full evaluation set are also reported. Main cross-seed results use means and seed ranges; per-run artifact tables retain Wilson 95% intervals under a binomial sampling interpretation for fixed checkpoints. Those intervals are descriptive for the finite CIFAR-10 test set and do not measure variability across training runs. Pairwise eligibility sets differ, so raw rates across model pairs are not strictly matched-population comparisons.

## 5 Experimental protocol

### 5.1 Data and model training

The pipeline uses the standard CIFAR-10 class order (Krizhevsky, 2009). A fixed split seed of 1729 partitions the 50,000 training images into 45,000 training and 5,000 validation images. Validation data receive no random augmentation. Training uses random crops with four-pixel padding and random horizontal flips. The 10,000 test images are reserved for final evaluation; checkpoint selection uses validation accuracy only.

All networks have ten output classes and are initialized without ImageNet weights. The ResNet-18 stem uses a  $3 \times 3$  convolution at stride one and no initial max pool. VGG16 uses its convolutional feature stack, global  $1 \times 1$  pooling, and a single 512-to-10 linear classifier. MobileNetV2 changes its initial convolution to stride one. These adaptations, especially the compact VGG classifier, must be stated when comparing with other implementations.

Inputs to the attack interface are floating-point pixels in  $[0, 1]$ . Channel normalization with mean  $(0.4914, 0.4822, 0.4465)$  and standard deviation  $(0.2470, 0.2435, 0.2616)$  occurs inside the model, so gradients include preprocessing while perturbation budgets remain in raw-pixel coordinates. Checkpoints contain architecture, class order, split indices, training seed, selected epoch, and model state. Evaluation rejects incompatible checkpoint metadata.

Training uses 200 epochs, SGD with momentum 0.9, weight decay  $5 \times 10^{-4}$ , cosine learning-rate decay, and batch size 128. ResNet-18 and MobileNetV2 use initial learning rate 0.1. That rate caused the VGG16 loss to remain near log 10 and validation accuracy near chance, so we stopped that failed run and selected 0.01 using validation performance only. We then froze 0.01 for all three VGG16 seeds before test evaluation. The independently trained seeds are 0, 1, and 2; the split seed remains 1729.

## 5.2 Attacks and controls

FGSM uses one signed cross-entropy gradient step with projection into the pixel range. PGD uses random initialization within the intersection of the pixel range and the  $L_\infty$  ball, 40 steps, absolute step size  $2/255$ , and five restarts at the primary budget  $8/255$ . Across iterates and restarts, the implementation prioritizes source misclassification, then larger source loss. Target predictions never influence this selection. This source-selected construction may differ from final-iterate PGD in its transfer behavior, so the selection rule is part of the attack specification.

The clean control leaves inputs unchanged. The random-noise control samples independent uniform perturbations within the same  $L_\infty$  budget and clips to the pixel range. The empirical scope is limited to the  $L_\infty$  threat model. We do not mix CW- $L_2$  results with these measurements or imply that equal numeric radii across norms represent equal threat models.

For every source and attack, the same generated inputs are evaluated on every target, including the source. Models remain in evaluation mode. Runtime checks reject nonfinite pixels, invalid pixel ranges, and norm violations. A saved example records its test index, true label, predictions, and actual perturbation. Visualization does not presume attack success or establish perceptual invisibility.

## 5.3 Reproducibility and sensitivity analysis

A seeded permutation specifies test indices, and the identical ordered sample list was used for every pair. All reported final evaluations used the full 10,000-image CIFAR-10 test set; the earlier 1,000-example runs were labeled as pilots and excluded from the empirical results. Each completed run saves arguments, software versions, commit identifier, checkpoint SHA-256 hashes, test indices, elapsed time, clean and attacked per-example predictions, and per-example  $L_2$  and  $L_\infty$  norms. Tables are generated from measured summaries rather than embedded constants.

We evaluate  $L_\infty$  radii  $2/255$ ,  $4/255$ , and  $8/255$ . The PGD step size is one quarter of the radius. At  $8/255$ , seed 0 uses 10, 40, and 100 steps with one restart and 100 steps with five restarts; seeds 1 and 2 repeat the 10-step, one-restart and 100-step, five-restart endpoints. At  $2/255$ , seed 0 also uses 100 steps and five restarts. Attack hyperparameters are not selected to maximize target transfer. We show training seeds individually and summarize seed-level estimates; we do not pool repeated test images as independent observations.

To test whether the architecture-specific VGG16 learning rate alone explains the largest asymmetries, we train seed-0 ResNet-18 and MobileNetV2 controls at 0.01 for the same 200 epochs and rerun the full-test  $8/255$  protocol with the existing VGG16 seed-0 checkpoint. This is a prespecified sensitivity comparison; it is not pooled with the three-seed primary estimates.

# 6 Results

## 6.1 Clean accuracy and primary attack strength

Table 1 reports validation-selected checkpoint performance on the untouched test set. Variation across the three training seeds is small relative to the differences among architectures. At the primary  $8/255$  budget, 40-step, five-restart PGD attains 100% source ASR for every ResNet-18 and MobileNetV2 checkpoint and at least 99.98% for every VGG16 checkpoint. The diagonal therefore confirms that the primary attack is strong against its source models.

Uniform random noise at the same  $8/255$  radius produces mean off-diagonal PTR between 1.45% and 7.89%, below the corresponding adversarial results. Noise CTR is based on much smaller source-success subsets and

Table 1: Clean test accuracy (%) by training seed. The final column is the mean and sample standard deviation over three checkpoints.

| Model       | Seed 0 | Seed 1 | Seed 2 | Mean $\pm$ SD    |
|-------------|--------|--------|--------|------------------|
| ResNet-18   | 95.11  | 95.01  | 95.20  | 95.11 $\pm$ 0.10 |
| VGG16       | 90.59  | 90.41  | 90.88  | 90.63 $\pm$ 0.24 |
| MobileNetV2 | 89.97  | 90.57  | 90.38  | 90.31 $\pm$ 0.31 |

Table 2: Mean source ASR (%) over three seeds.

| Source      | FGSM 8/255 | PGD 2/255 | PGD 4/255 | PGD 8/255 |
|-------------|------------|-----------|-----------|-----------|
| ResNet-18   | 58.54      | 94.47     | 99.91     | 100.00    |
| VGG16       | 87.31      | 71.73     | 97.93     | 99.98     |
| MobileNetV2 | 92.03      | 99.85     | 100.00    | 100.00    |

is not used to rank model pairs. Because each nominal source row uses a separately seeded noise realization, these controls establish a non-adversarial baseline within each row rather than a paired comparison among noise rows.

## 6.2 Directionality at the primary budget

Table 3 shows off-diagonal PGD PTR at 8/255. Transfer is strongly directional. VGG16 examples transfer to ResNet-18 at 92.73% on average, whereas the reverse direction reaches 18.09%. MobileNetV2-to-VGG16 is 25.94%, whereas VGG16-to-MobileNetV2 is 95.23%. The ordering persists across all three seeds, so it is not produced by one checkpoint draw. These are associations among the trained architectures and recipes; differences in capacity, optimization, and clean accuracy preclude an architecture-only causal interpretation.

The matched-learning-rate control preserves the two largest VGG asymmetries (Table 4). VGG16-to-ResNet-18 exceeds the reverse direction by 63.27 points, and VGG16-to-MobileNetV2 exceeds the reverse by 74.57 points. The control ResNet-18 has 93.82% clean test accuracy, compared with 95.11% for its primary seed-0 checkpoint; control MobileNetV2 has 90.90%, compared with 89.97%. Individual transfer entries change by up to 25.93 points across the full matrix. The control rejects a learning-rate-only account of the VGG asymmetry; it licenses no stronger causal claim because the resulting checkpoints also differ in accuracy, optimization trajectory, and loss geometry.

## 6.3 Budget and conditioning sensitivity

Transfer falls sharply with the perturbation budget, but the directional ordering is broadly stable (Figure 1). VGG16-to-ResNet-18 drops from 92.73% at 8/255 to 15.45% at 2/255; ResNet-18-to-MobileNetV2 drops from 72.27% to 21.61%. These changes rule out treating one reported transfer percentage as a property of a model pair independent of its threat model.

At 2/255, VGG16 source ASR is only 71.05–72.98% (Table 2). Table 6 applies the exact conditioning identity with absolute counts. Conditioning is nearly mechanical in five directions: mean  $b_{st}$  is at most 0.68%, and mean CTR is within 0.27 points of the value obtained by dividing mean PTR by mean  $a_{st}$ . Across individual seeds, the largest corresponding deviation is 0.40 points. VGG16-to-MobileNetV2 is the exception, with mean  $b_{st}$  = 2.83% and observed CTR 1.19 points below the zero-residual value. This exception confirms

Table 3: PGD PTR (%) at 8/255, reported as mean [minimum, maximum] over three training seeds.

| Source      | ResNet-18           | VGG16               | MobileNetV2         |
|-------------|---------------------|---------------------|---------------------|
| ResNet-18   | —                   | 18.09 [16.73,19.38] | 72.27 [71.76,72.73] |
| VGG16       | 92.73 [92.40,92.99] | —                   | 95.23 [94.80,95.61] |
| MobileNetV2 | 67.60 [63.46,70.61] | 25.94 [25.73,26.29] | —                   |

Table 4: Seed-0 PGD PTR (%) at 8/255 before and after training all three architectures with initial learning rate 0.01.

| Direction                       | Primary | Matched LR |
|---------------------------------|---------|------------|
| ResNet-18 $\rightarrow$ VGG16   | 19.38   | 32.71      |
| VGG16 $\rightarrow$ ResNet-18   | 92.99   | 95.98      |
| MobileNetV2 $\rightarrow$ VGG16 | 25.73   | 18.55      |
| VGG16 $\rightarrow$ MobileNetV2 | 94.80   | 93.12      |

that the successful- and failed-source strata are not interchangeable. More fundamentally, conventional source ASR and PTR cannot determine CTR even when the residual is negligible, because conventional ASR uses a broader, target-independent denominator than  $a_{st}$ .

The stronger 100-step, five-restart 8/255 endpoint changes off-diagonal PTR by at most 1.70 percentage points relative to the primary 40-step, five-restart setting across all seeds. Increasing optimization does not increase every transfer entry: for example, ResNet-18-to-VGG16 falls slightly. At 2/255 on seed 0, increasing from 40 to 100 steps at five restarts changes all off-diagonal PTR entries by at most 0.27 points and CTR entries by at most 0.40 points. This supports convergence of the primary step count without asserting a global optimum.

FGSM reinforces the distinction between source strength and transfer (Table 5). At 8/255, its mean PTR is 27.80% for ResNet-18-to-VGG16 and 29.34% for MobileNetV2-to-VGG16, compared with 18.09% and 25.94% for PGD. PGD nevertheless has far higher source ASR. For the other four off-diagonal directions, PGD transfers more strongly than FGSM. This replicates prior observations that iterative source optimization can reduce transfer (Dong et al., 2018); an attack ranking depends on direction and cannot be inferred from white-box source success alone.


## 7 Evidence status and limitations

An anonymized reproducibility package accompanies this submission. It contains the corrected source snapshot, aggregate tables, manifests, and per-example evaluation records; the public repository and archival identifier are withheld during double-blind review. The final study record contains nine checkpoints, full-test predictions, exact eligibility counts, perturbation norms, manifests, and checkpoint and source hashes. The three seed triads provide evidence about variability under the specified training procedure, but three draws support descriptive ranges rather than precise population-level inference. Per-run Wilson intervals describe fixed-checkpoint binomial variation and do not capture training variability. The same 10,000 test images recur across seeds, so image-level counts must not be pooled as 30,000 independent observations.

The study covers three standard CNN families on CIFAR-10 and untargeted  $L_\infty$  attacks. It does not establish behavior on larger datasets, transformers, adversarially trained models, targeted attacks, or physical settings. It also does not compare against every modern transfer attack. The VGG16 learning-rate adjustment was selected using validation results after the common recipe failed. A seed-0 matched-learning-rate control preserves its strongest asymmetries, but only one control seed was run and some individual entries change substantially; optimization and architecture remain entangled. Random-noise draws are independently generated for each nominal source row, so noise-control rows are reproducible but do not share one

Table 5: FGSM PTR (%) at 8/255, reported as mean [minimum, maximum] over three seeds.

| Source      | ResNet-18           | VGG16               | MobileNetV2         |
|-------------|---------------------|---------------------|---------------------|
| ResNet-18   | —                   | 27.80 [26.21,30.01] | 58.30 [56.73,59.43] |
| VGG16       | 55.27 [54.14,56.75] | —                   | 64.88 [64.38,65.71] |
| MobileNetV2 | 47.37 [47.20,47.55] | 29.34 [28.09,30.55] | —                   |

Figure 1: Mean PGD PTR across three seeds versus  $L_\infty$  budget. Error bars span the minimum and maximum seed values.

common perturbation realization. The retained manifests record a dirty working tree at execution time; they preserve prediction checksums, checkpoint hashes, the recorded commit, and hashes of the executed source files, while the clean release tag archives the final corrected code.

The original artifact is a motivating self-correction, not evidence that similar errors are common elsewhere. The present results do not imply that one architecture is intrinsically more transferable: architecture, parameterization, training dynamics, and clean accuracy vary together. PTR and CTR answer different questions, and neither should be selected after seeing which produces the more dramatic result.

Table 6: 2/255 PGD conditioning decomposition. Count tuples list seeds 0/1/2; rates are mean [range] percentages.

| Direction | Pairwise eligible $ E $ | Source success $ F $ |
|-----------|-------------------------|----------------------|
| R→V       | 8905/8874/8923          | 8381/8318/8424       |
| R→M       | 8831/8894/8873          | 8307/8339/8374       |
| V→R       | 8905/8874/8923          | 6457/6257/6303       |
| V→M       | 8583/8589/8620          | 6139/5976/6003       |
| M→R       | 8831/8894/8873          | 8822/8879/8857       |
| M→V       | 8583/8589/8620          | 8574/8574/8604       |

---

| Direction | Pair ASR            | Fail-target      | PTR   | CTR   |
|-----------|---------------------|------------------|-------|-------|
| R→V       | 94.09 [93.73,94.41] | 0.00 [0.00,0.00] | 5.84  | 6.21  |
| R→M       | 94.07 [93.76,94.38] | 0.25 [0.19,0.36] | 21.61 | 22.96 |
| V→R       | 71.22 [70.51,72.51] | 0.68 [0.42,0.96] | 15.45 | 21.42 |
| V→M       | 70.25 [69.58,71.53] | 2.83 [2.71,2.99] | 23.29 | 31.97 |
| M→R       | 99.85 [99.82,99.90] | 0.00 [0.00,0.00] | 10.79 | 10.81 |
| M→V       | 99.84 [99.81,99.90] | 0.00 [0.00,0.00] | 4.61  | 4.62  |

## 8 Broader impact

This work studies standard, bounded attacks on CIFAR-10 classifiers to improve the validity and reproducibility of adversarial-robustness evaluation. The same methods can inform attempts to evade classifiers, but the experiments do not target a deployed system, use private data, or demonstrate a physical-world attack. Releasing exact denominators, perturbation norms, controls, and evaluation code primarily supports defensive auditing by making misleading transfer claims easier to detect and reproduce.

## 9 Conclusion

Across three training seeds, adversarial transfer on CIFAR-10 is strongly directional and budget dependent. At 8/255, mean PGD PTR spans 18.09–95.23% across off-diagonal pairs; at 2/255, it spans 4.61–23.29%. Conditioning is nearly mechanical in five directions, with one measurable failed-source exception. Yet conventional source ASR and PTR still cannot reconstruct CTR because source ASR uses the wrong, target-independent denominator. The strongest VGG directionality persists in a seed-0 matched-learning-rate control, while the movement of individual entries cautions against an architecture-only interpretation. Stronger source optimization can also reduce transfer for a particular direction, replicating established source-overfitting behavior. Compatible labels and preprocessing make these measurements meaningful; clean-error exclusions, explicit denominators, seed-level reporting, and convergence checks make their limits visible.

## References

- Nicholas Carlini and David Wagner. Towards evaluating the robustness of neural networks. In *IEEE Symposium on Security and Privacy*, 2017. URL <https://arxiv.org/abs/1608.04644>.
- Nicholas Carlini, Anish Athalye, Nicolas Papernot, Wieland Brendel, Jonas Rauber, Dimitris Tsipras, Ian Goodfellow, Aleksander Madry, and Alexey Kurakin. On evaluating adversarial robustness. *arXiv preprint arXiv:1902.06705*, 2019. URL <https://arxiv.org/abs/1902.06705>.
- Francesco Croce and Matthias Hein. Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pp. 2206–2216, 2020. URL <https://proceedings.mlr.press/v119/croce20b.html>.
- Francesco Croce, Maksym Andriushchenko, Vikash Sehwal, Edoardo Debenedetti, Nicolas Flammarion, Mung Chiang, Prateek Mittal, and Matthias Hein. Robustbench: A standardized adversarial robustness benchmark. In *NeurIPS Datasets and Benchmarks Track*, 2021. URL <https://robustbench.github.io/>.
- Ambra Demontis, Marco Melis, Maura Pintor, Matthew Jagielski, Battista Biggio, Alina Oprea, Cristina Nita-Rotaru, and Fabio Roli. Why do adversarial attacks transfer? explaining transferability of evasion and poisoning attacks. In *28th USENIX Security Symposium*, 2019. URL <https://www.usenix.org/conference/usenixsecurity19/presentation/demontis>.
- Yinpeng Dong, Fangzhou Liao, Tianyu Pang, Hang Su, Jun Zhu, Xiaolin Hu, and Jianguo Li. Boosting adversarial attacks with momentum. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2018. URL [https://openaccess.thecvf.com/content\\_cvpr\\_2018/html/Dong\\_Boosting\\_Adversarial\\_Attacks\\_CVPR\\_2018\\_paper.html](https://openaccess.thecvf.com/content_cvpr_2018/html/Dong_Boosting_Adversarial_Attacks_CVPR_2018_paper.html).
- Ian J. Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial examples. In *International Conference on Learning Representations*, 2015. URL <https://arxiv.org/abs/1412.6572>.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009. URL <https://www.cs.toronto.edu/~kriz/cifar.html>.
- Yanpei Liu, Xinyun Chen, Chang Liu, and Dawn Song. Delving into transferable adversarial examples and black-box attacks. In *International Conference on Learning Representations*, 2017. URL <https://arxiv.org/abs/1611.02770>.
- Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu. Towards deep learning models resistant to adversarial attacks. In *International Conference on Learning Representations*, 2018. URL <https://arxiv.org/abs/1706.06083>.
- Nicolas Papernot, Patrick McDaniel, Ian Goodfellow, Somesh Jha, Z. Berkay Celik, and Ananthram Swami. Practical black-box attacks against machine learning. In *ACM Asia Conference on Computer and Communications Security*, 2017. URL <https://arxiv.org/abs/1602.02697>.
- Florian Tramèr, Nicolas Papernot, Ian Goodfellow, Dan Boneh, and Patrick McDaniel. The space of transferable adversarial examples. *arXiv preprint arXiv:1704.03453*, 2017. URL <https://arxiv.org/abs/1704.03453>.
- Florian Tramèr, Nicholas Carlini, Wieland Brendel, and Aleksander Madry. On adaptive attacks to adversarial example defenses. In *Advances in Neural Information Processing Systems*, 2020. URL <https://arxiv.org/abs/2002.08347>.